Global Enterprise AI Governance and Acceptable Use Policy

Policy ID: POL-SEC-2025-04 · Effective Date: January 1, 2025

1.0 Purpose and Scope

This policy establishes governance boundaries, data protection mandates, and security requirements governing the deployment and utilization of Large Language Model (LLM) services and Retrieval-Augmented Generation (RAG) applications across corporate environments.

2.0 Mandatory Security Controls

2.1 Zero Persistent Secret Logging: API credentials, bearer tokens, and cryptographic keys must never be persisted to file logs, client browsers, or unencrypted storage. Ephemeral in-memory handling with explicit session termination is mandatory.

2.2 Document Data Isolation: User documents processed within temporary sessions must be isolated per tenant and purged upon session disconnect. No document data may be utilized for model pre-training without explicit customer written consent.

2.3 Evidence and Attribution Requirements: All AI-synthesized responses regarding contractual, technical, or financial facts must provide verifiable, page-indexed citations to source material.